

# Breast Cancer Prediction Case Study - Bayesian Logistic Regression with Comparison of Frequentist and Bayesian Variable Selection Methods

Laura Silvana Alvarez Luque, Florencia Luque, Nicolas Bühringer, Simon Schmetz

2025-03-14

## Introduction

This case study is based on the Breast Cancer Wisconsin (Diagnostic) Data Set (<https://www.kaggle.com/datasets/uciml/breast-cancer-wisconsin-data>). The data set contains 569 observations and 32 variables. The data set is available at the UCI Machine Learning Repository. The data set contains mean (and at times min and max) values of the patient for the following numeric (continuous) variables:

- a) radius (mean of distances from center to points on the perimeter)
- b) texture (standard deviation of gray-scale values)
- c) perimeter
- d) area
- e) smoothness (local variation in radius lengths)
- f) compactness ( $\text{perimeter}^2 / \text{area} - 1.0$ )
- g) concavity (severity of concave portions of the contour)
- h) concave points (number of concave portions of the contour)
- i) symmetry
- j) fractal dimension ("coastline approximation" - 1)

The data set also contains the following Binary variables:

- 2) Diagnosis (M = malignant, B = benign)

Where Malignant (M) means the tumor is cancerous, while Benign (B): means that the tumor is non-cancerous.

## Read Data

```
data <- read.csv("data.csv", header = TRUE, sep = ",")
data <- dplyr::select(data, -c(X,id))
names(data) <- gsub("\\.", "_", names(data))
data$diagnosis <- ifelse(data$diagnosis == "M", 1, 0)
```

## Exploratory Data Analysis

### Relation with response var

```
numeric_vars <- data %>%select_if(is.numeric) %>% colnames()
numeric_vars <- setdiff(numeric_vars, "diagnosis")
```

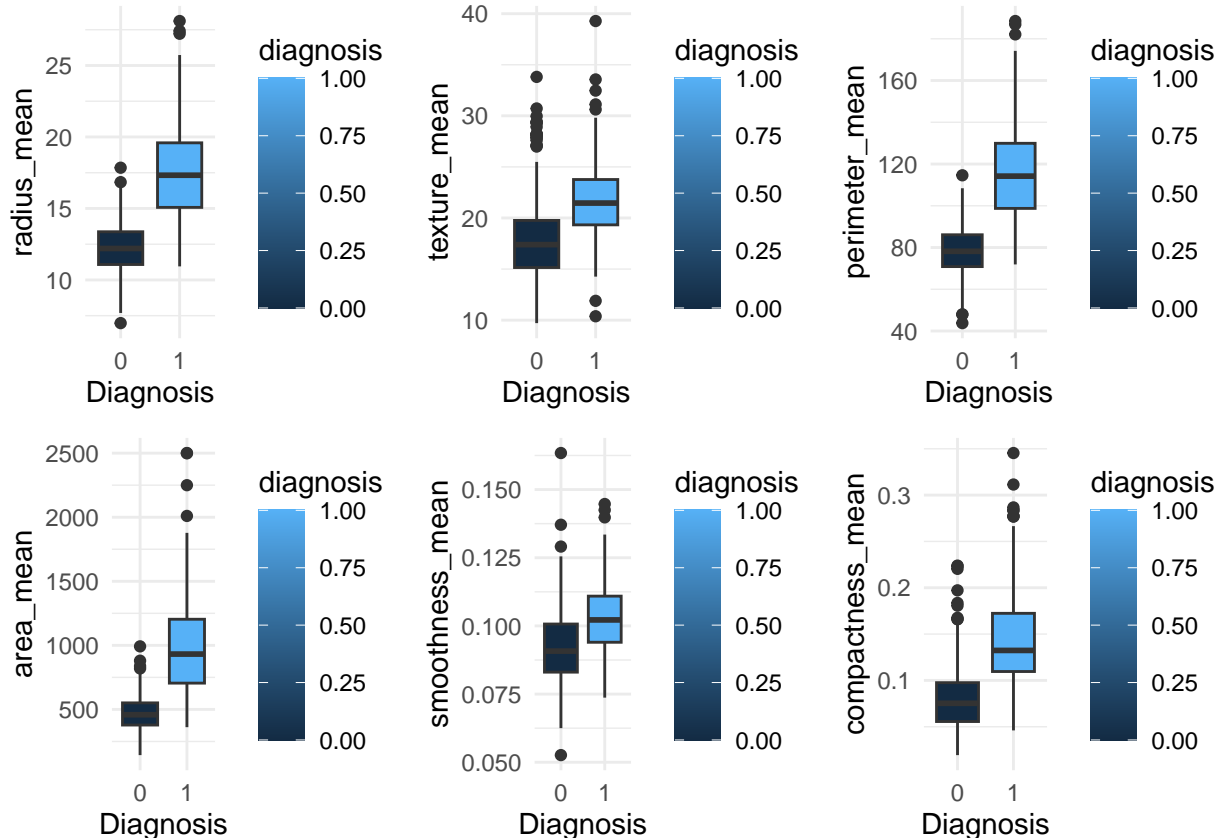
```
plots <- lapply(numeric_vars, function(var) {
  ggplot(data, aes(x = factor(diagnosis), y = .data[[var]], fill = diagnosis)) +
    geom_boxplot() +
    labs(x = "Diagnosis", y = var) +
    theme_minimal()
})
```

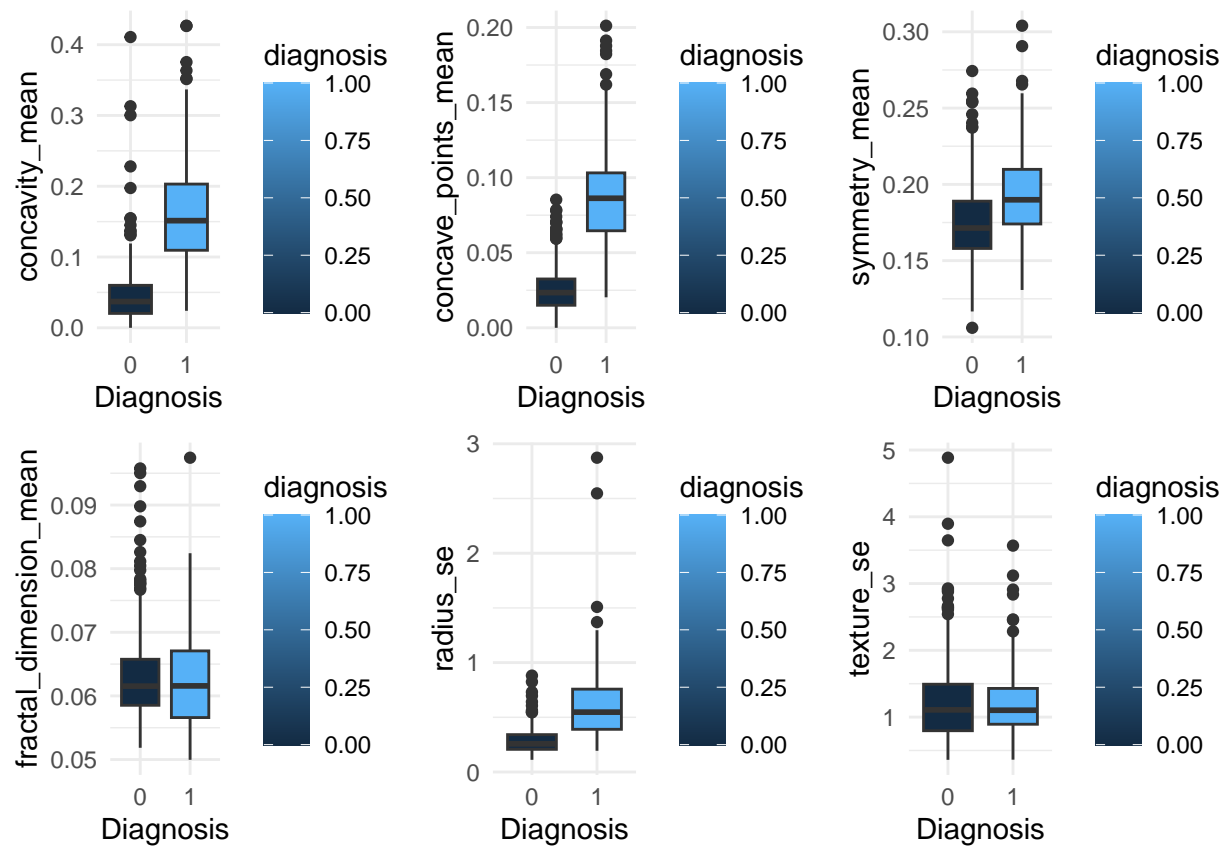
```
# Print all plots
# install.packages("gridExtra")
library(gridExtra)
```

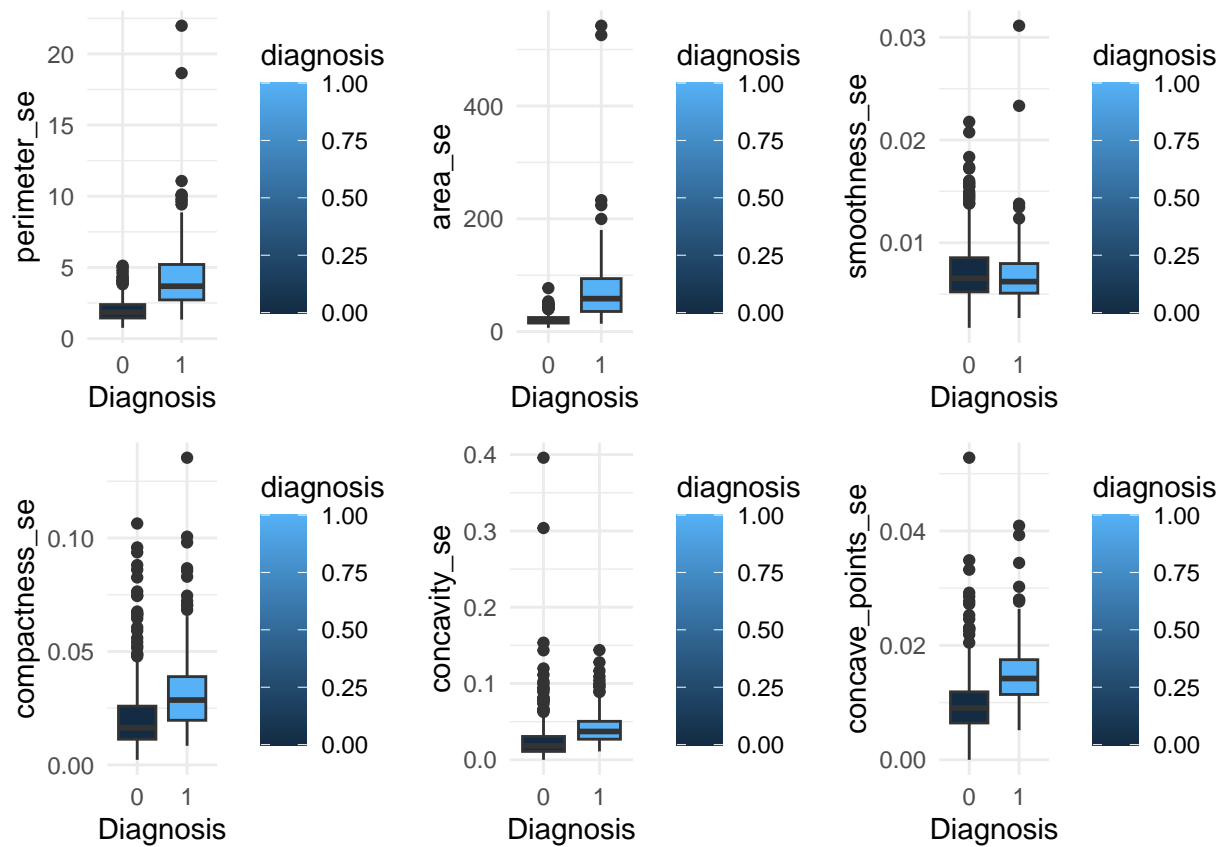
```
##
## Attaching package: 'gridExtra'
## The following object is masked from 'package:dplyr':
##
## combine
```

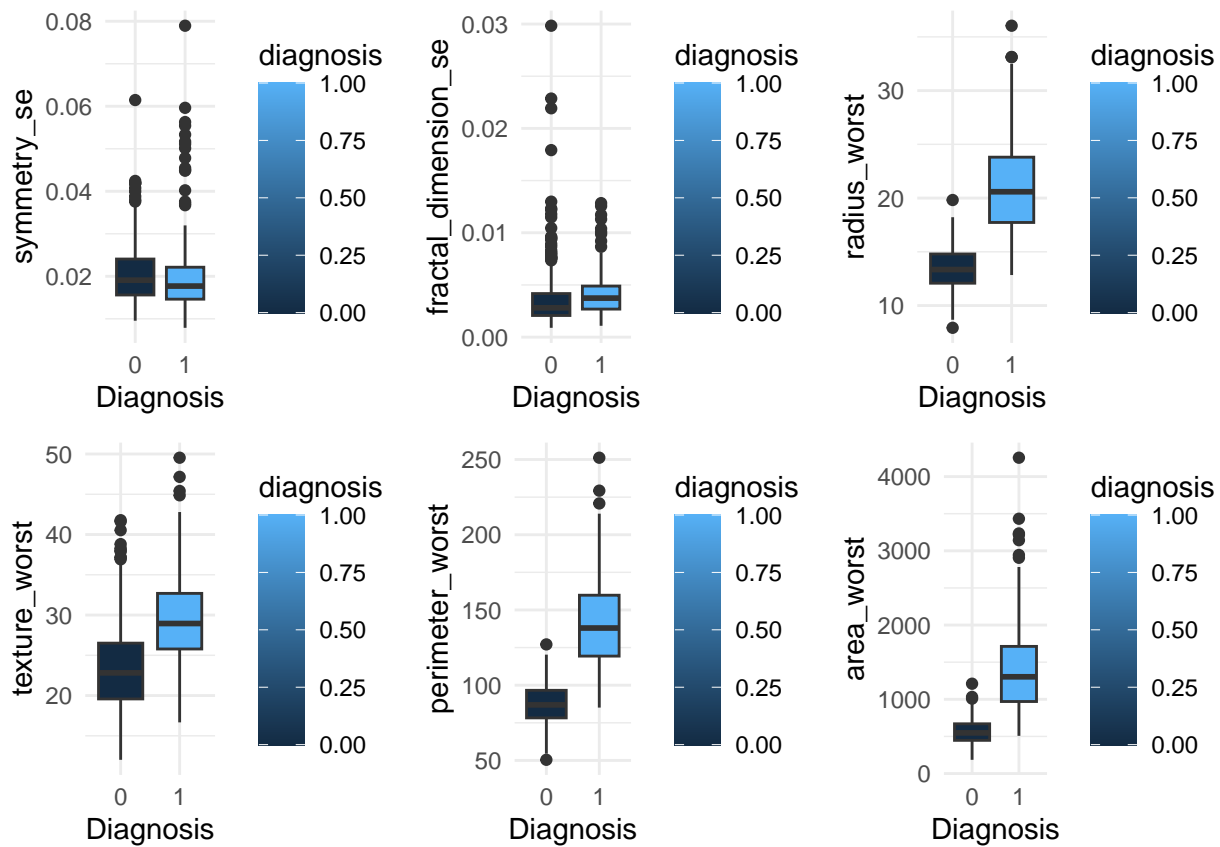
```
# Display plots in batches of 6 (2 rows x 3 columns)
num_plots <- length(plots)
batch_size <- 6

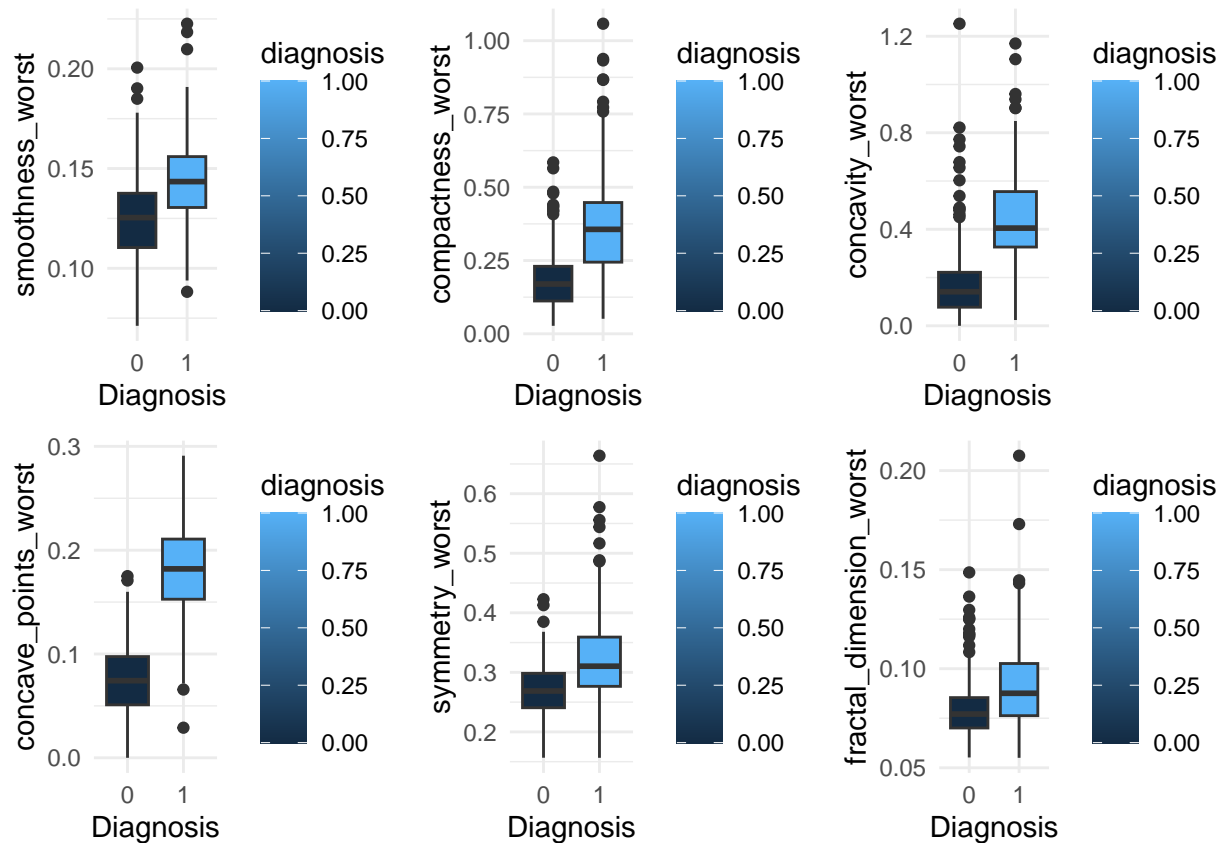
for(i in seq(1, num_plots, batch_size)) {
  end_idx <- min(i + batch_size - 1, num_plots)
  batch_plots <- plots[i:end_idx]
  grid.arrange(grobs = batch_plots, ncol = 3)
}
```











```
summary_stat <- data %>%
  group_by(factor(diagnosis)) %>%
  summarise(across(all_of(numeric_vars),
    list(
      mean = ~mean(.x, na.rm = TRUE),
      sd = ~sd(.x, na.rm = TRUE),
      median = ~median(.x, na.rm = TRUE),
      min = ~min(.x, na.rm = TRUE),
      max = ~max(.x, na.rm = TRUE)
    )
  ))
```

*# For easier viewing, you can pivot longer*

```
summary_long <- summary_stat %>%
  pivot_longer(cols = ~factor(diagnosis)`,
    names_to = c("variable", "stat"),
    names_pattern = "(.*)_(.*)")
```

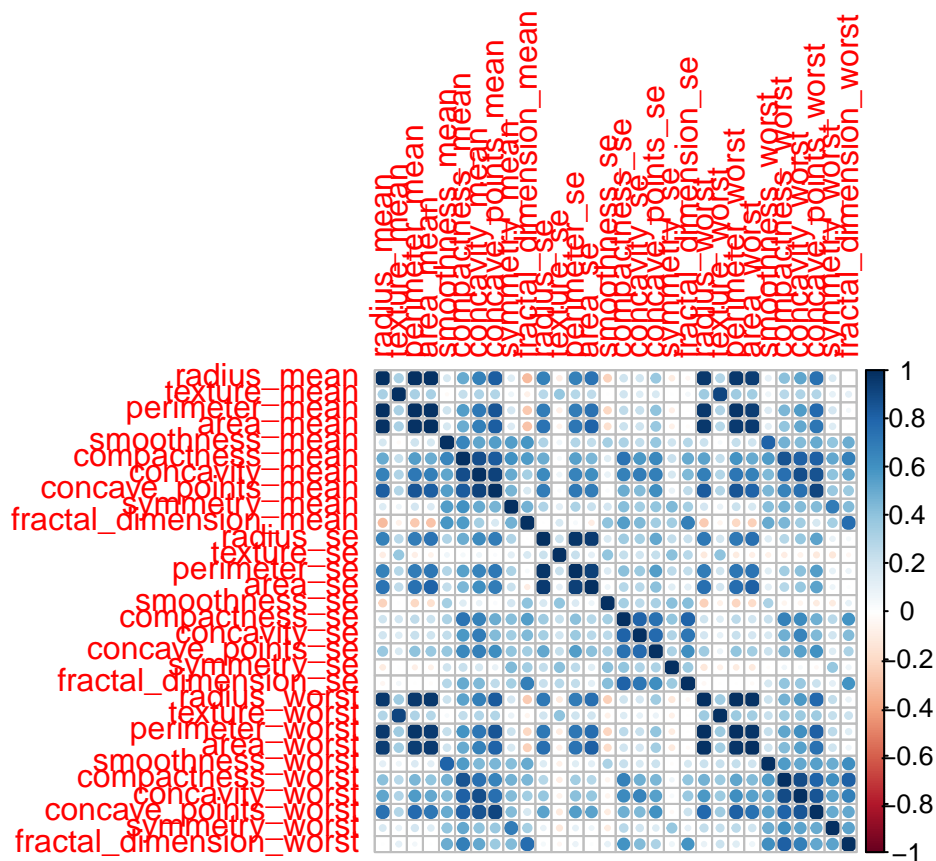
```
summary_long
```

```
## # A tibble: 300 x 4
##   `factor(diagnosis)` variable    stat    value
##   <fct>                <chr>      <chr> <dbl>
## 1 0                    radius_mean mean    12.1
## 2 0                    radius_mean sd      1.78
```

```
## 3 0          radius_mean median 12.2
## 4 0          radius_mean min    6.98
## 5 0          radius_mean max    17.8
## 6 0          texture_mean mean   17.9
## 7 0          texture_mean sd     4.00
## 8 0          texture_mean median 17.4
## 9 0          texture_mean min     9.71
## 10 0         texture_mean max    33.8
## # i 290 more rows
```

## Correlation

```
# Check correlation between numeric variables
cor_matrix <- cor(data[, numeric_vars])
corrplot(cor_matrix, method = "circle")
```



```
# Or find highly correlated variables
high_cor <- findCorrelation(cor_matrix, cutoff = 0.8)
problematic_vars <- numeric_vars[high_cor]
print(problematic_vars)
```

```
## [1] "concavity_mean"      "concave_points_mean" "compactness_mean"
## [4] "concave_points_worst" "concavity_worst"     "perimeter_worst"
## [7] "radius_worst"        "perimeter_mean"     "compactness_worst"
## [10] "area_worst"          "radius_mean"        "perimeter_se"
## [13] "compactness_se"      "area_se"            "smoothness_mean"
## [16] "texture_mean"
```

# Variable Selection

## Frequentist Approach

Check VIF and remove variables with extremely high values

```
predictors <- setdiff(names(data), c("diagnosis"))
formula_str <- paste("diagnosis ~", paste(predictors, collapse = " + "))
formula <- as.formula(formula_str)

l_reg = lm(formula, data)
vif_values <- vif(l_reg)

vif_df <- data.frame(
  Variable = names(vif_values),
  VIF = vif_values
)
vif_df <- vif_df %>% arrange(desc(VIF))
print(vif_df)
```

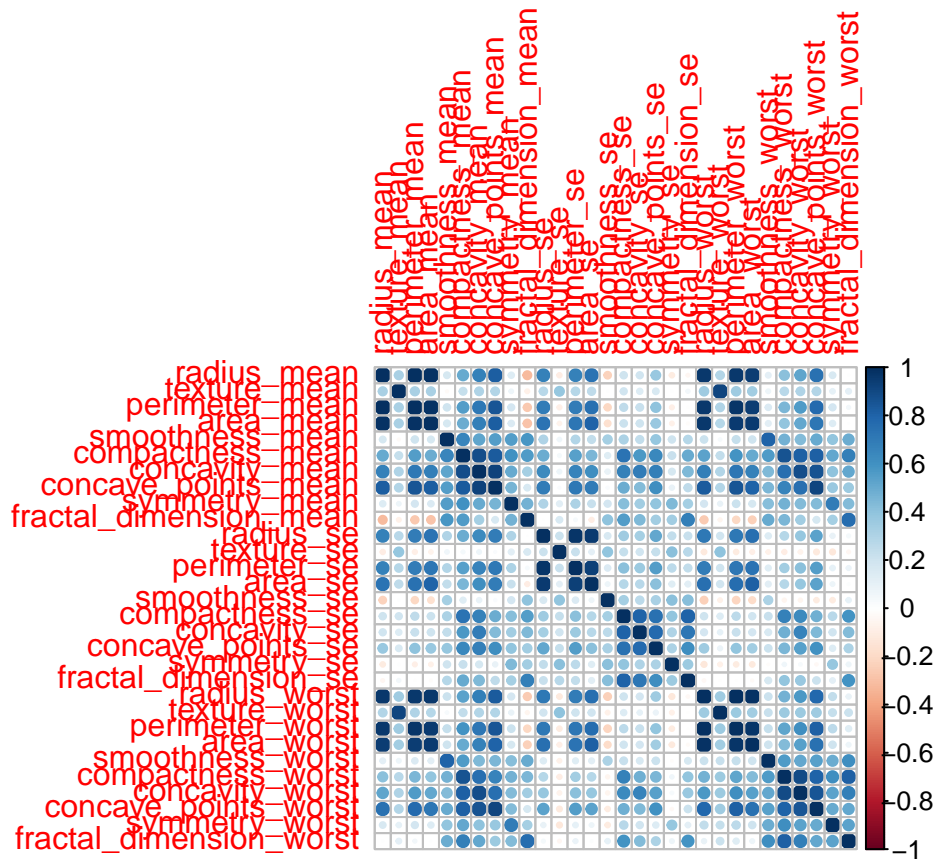
##	Variable	VIF
## radius_mean	radius_mean	3806.115296
## perimeter_mean	perimeter_mean	3786.400419
## radius_worst	radius_worst	799.105946
## perimeter_worst	perimeter_worst	405.023336
## area_mean	area_mean	347.878657
## area_worst	area_worst	337.221924
## radius_se	radius_se	75.462027
## concavity_mean	concavity_mean	70.767720
## perimeter_se	perimeter_se	70.359695
## concave_points_mean	concave_points_mean	60.041733
## compactness_mean	compactness_mean	50.505168
## area_se	area_se	41.163091
## compactness_worst	compactness_worst	36.982755
## concave_points_worst	concave_points_worst	36.763714
## concavity_worst	concavity_worst	31.970723
## fractal_dimension_worst	fractal_dimension_worst	18.861533
## texture_worst	texture_worst	18.569966
## fractal_dimension_mean	fractal_dimension_mean	15.756977
## concavity_se	concavity_se	15.694833
## compactness_se	compactness_se	15.366324
## texture_mean	texture_mean	11.884048
## concave_points_se	concave_points_se	11.520796
## smoothness_worst	smoothness_worst	10.923061
## fractal_dimension_se	fractal_dimension_se	9.717987
## symmetry_worst	symmetry_worst	9.520570
## smoothness_mean	smoothness_mean	8.194282
## symmetry_se	symmetry_se	5.175426
## symmetry_mean	symmetry_mean	4.220656
## texture_se	texture_se	4.205423
## smoothness_se	smoothness_se	4.027923

```
vars_to_exclude <- c(head(vif_df,15)$Variable)
# setdiff(problematic_vars , vars_to_exclude)
# setdiff(vars_to_exclude, problematic_vars)
```



Check correlations after excluding x VIF, variables to pay attention if something does not work.

```
# Check correlation between numeric variables
cor_matrix_f <- cor(data[, setdiff(numeric_vars, vars_to_exclude)])
corrplot(cor_matrix, method = "circle")
```



```
# Or find highly correlated variables
high_cor_f <- findCorrelation(cor_matrix_f, cutoff = 0.8)
problematic_vars_f <- setdiff(numeric_vars, vars_to_exclude)[high_cor_f]
print(problematic_vars_f)
```

```
## [1] "compactness_se" "smoothness_mean" "texture_mean"
selected_freq <- setdiff(numeric_vars, vars_to_exclude)
```

## Bayesian Approach

```
library(BAS)

# Fit a Bayesian logistic regression with variable selection
model_bas <- bas.glm(diagnosis ~ .,
  data = data,
  family = binomial(),
  method = "MCMC", # or "BAS" for deterministic sampling
  MCMC.iterations = 10000,
  modelprior = uniform()) # Prior over model space
```

```
# Summary of results
summary(model_bas)
```

##	P(B != 0   Y)	model 1	model 2	model 3
## Intercept	1.0000	1.00000000	1.00000	1.00000000
## radius_mean	0.3486	0.00000000	0.00000	1.00000000
## texture_mean	0.2364	0.00000000	0.00000	0.00000000
## perimeter_mean	0.2586	0.00000000	0.00000	0.00000000
## area_mean	0.2374	0.00000000	1.00000	0.00000000
## smoothness_mean	0.2168	0.00000000	0.00000	0.00000000
## compactness_mean	0.4635	0.00000000	1.00000	0.00000000
## concavity_mean	0.3506	0.00000000	0.00000	1.00000000
## concave_points_mean	0.5527	0.00000000	1.00000	0.00000000
## symmetry_mean	0.1601	0.00000000	0.00000	0.00000000
## fractal_dimension_mean	0.1654	0.00000000	0.00000	0.00000000
## radius_se	0.4634	1.00000000	0.00000	0.00000000
## texture_se	0.3690	0.00000000	0.00000	0.00000000
## perimeter_se	0.2660	0.00000000	0.00000	0.00000000
## area_se	0.5434	0.00000000	1.00000	1.00000000
## smoothness_se	0.3558	1.00000000	0.00000	0.00000000
## compactness_se	0.4479	0.00000000	0.00000	0.00000000
## concavity_se	0.2973	0.00000000	0.00000	0.00000000
## concave_points_se	0.4684	0.00000000	1.00000	1.00000000
## symmetry_se	0.2091	0.00000000	0.00000	0.00000000
## fractal_dimension_se	0.5357	1.00000000	1.00000	1.00000000
## radius_worst	0.2816	0.00000000	0.00000	1.00000000
## texture_worst	0.8778	1.00000000	1.00000	1.00000000
## perimeter_worst	0.2354	0.00000000	0.00000	0.00000000
## area_worst	0.6316	1.00000000	1.00000	0.00000000
## smoothness_worst	0.4006	0.00000000	0.00000	1.00000000
## compactness_worst	0.2426	1.00000000	0.00000	0.00000000
## concavity_worst	0.4138	1.00000000	0.00000	0.00000000
## concave_points_worst	0.4240	1.00000000	0.00000	0.00000000
## symmetry_worst	0.5546	0.00000000	1.00000	0.00000000
## fractal_dimension_worst	0.4300	1.00000000	1.00000	1.00000000
## BF	NA	0.03514088	1.00000	0.06543644
## PostProbs	NA	0.00500000	0.00460	0.00270000
## R2	NA	0.92380000	0.93790	0.92540000
## dim	NA	10.00000000	11.00000	10.00000000
## logmarg	NA	-52.96576776	-49.61738	-52.34405369
##	model 4	model 5		
## Intercept	1.00000000	1.0000000		
## radius_mean	1.00000000	0.0000000		
## texture_mean	0.00000000	1.0000000		
## perimeter_mean	0.00000000	0.0000000		
## area_mean	0.00000000	0.0000000		
## smoothness_mean	0.00000000	0.0000000		
## compactness_mean	1.00000000	1.0000000		
## concavity_mean	0.00000000	0.0000000		
## concave_points_mean	1.00000000	1.0000000		
## symmetry_mean	0.00000000	0.0000000		
## fractal_dimension_mean	0.00000000	0.0000000		
## radius_se	0.00000000	0.0000000		
## texture_se	0.00000000	0.0000000		

```
## perimeter_se      1.00000000  0.00000000
## area_se           1.00000000  1.00000000
## smoothness_se     0.00000000  0.00000000
## compactness_se    0.00000000  0.00000000
## concavity_se       0.00000000  0.00000000
## concave_points_se  1.00000000  0.00000000
## symmetry_se        0.00000000  1.00000000
## fractal_dimension_se 1.00000000  1.00000000
## radius_worst       0.00000000  1.00000000
## texture_worst      1.00000000  0.00000000
## perimeter_worst    1.00000000  0.00000000
## area_worst         0.00000000  0.00000000
## smoothness_worst   0.00000000  1.00000000
## compactness_worst  1.00000000  0.00000000
## concavity_worst    0.00000000  1.00000000
## concave_points_worst 0.00000000  0.00000000
## symmetry_worst     1.00000000  1.00000000
## fractal_dimension_worst 1.00000000  0.00000000
## BF                 0.06601369  0.0113394
## PostProbs          0.00240000  0.0022000
## R2                 0.94130000  0.9263000
## dim                13.00000000  11.0000000
## logmarg            -52.33527083 -54.0968492
```

```
# Posterior inclusion probabilities
pip <- model_bas$probne0
variable_names <- names(pip)
#pip_df <- data.frame(Variable = numeric_vars,
#                      InclusionProb = pip)
#pip_df <- pip_df[order(pip_df$InclusionProb, decreasing = TRUE),]
#print(pip_df)

selected_bayes <- c( "perimeter_mean", "concave_points_mean", "compactness_mean",
                    "concavity_mean", "area_se", "smoothness_se", "concave_points_se",
                    "fractal_dimension_se", "radius_worst", "texture_worst",
                    "fractal_dimension_worst")
```

## Logistic Models

### Freq var selection

```
formula_str <- paste("diagnosis ~", paste(selected_freq, collapse = " + "))
formula <- as.formula(formula_str)

freq_model1 <- lm(formula, data = data)
beta.start1 <- coef(freq_model1)

out = MCMClogit(formula, data, burnin=1000, mcmc=21000, beta.start = beta.start1)
summary(out)

##
## Iterations = 1001:22000
## Thinning interval = 1
## Number of chains = 1
```

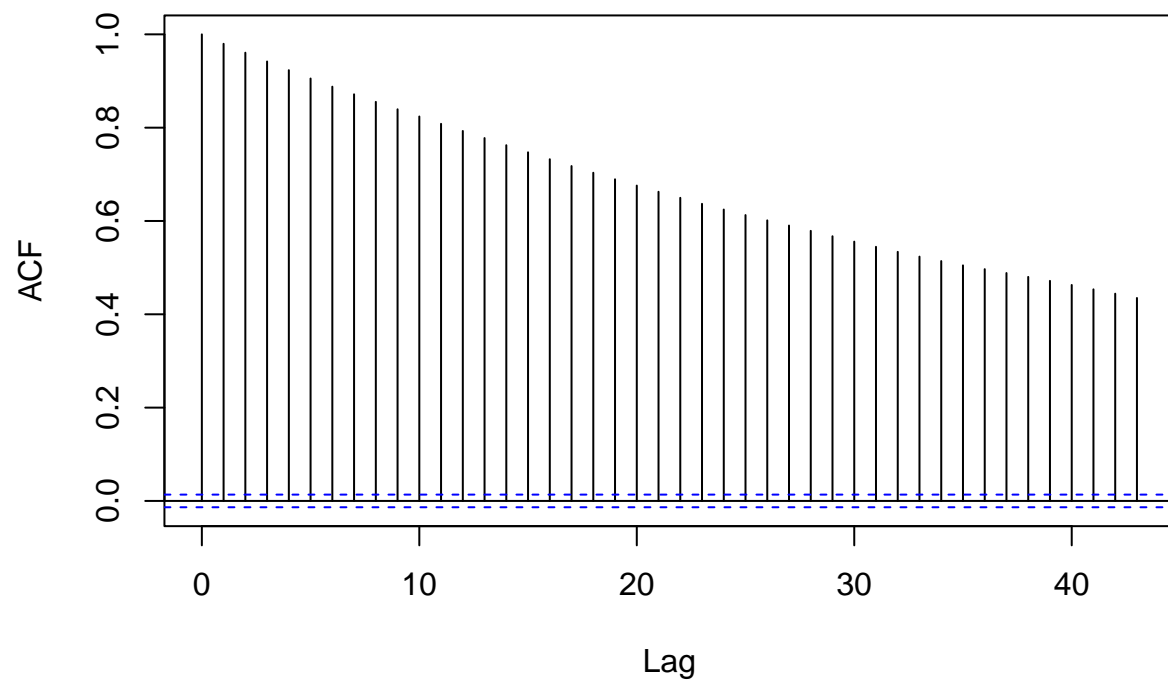
```

## Sample size per chain = 21000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##              Mean      SD Naive SE Time-series SE
## (Intercept)   -7.0448   2.7936 0.0192780      0.194235
## texture_mean    0.1402   0.1332 0.0009190      0.009217
## smoothness_mean 212.5324 39.8023 0.2746618      2.930954
## symmetry_mean    6.1214 13.6580 0.0942492      1.029750
## fractal_dimension_mean -692.6680 86.5229 0.5970649      6.070409
## texture_se     -1.4180   0.7611 0.0052523      0.051882
## smoothness_se  -244.3942 113.0283 0.7799695      7.684728
## compactness_se   22.7986 27.5534 0.1901367      2.002600
## concavity_se     -1.5990 12.5088 0.0863193      0.973608
## concave_points_se 419.3653 67.1342 0.4632698      4.945145
## symmetry_se       0.1725 46.5293 0.3210826      3.015775
## fractal_dimension_se -254.7769 237.3704 1.6380116     16.019617
## texture_worst     0.2644   0.1237 0.0008536      0.008767
## smoothness_worst  -0.7513 26.7957 0.1849080      1.771922
## symmetry_worst    12.2862   8.1680 0.0563646      0.596335
## fractal_dimension_worst 155.0970 37.1387 0.2562813      2.341975
##
## 2. Quantiles for each variable:
##
##              2.5%      25%      50%      75%      97.5%
## (Intercept)  -12.45556  -8.79360  -7.0801  -5.1257  -1.54107
## texture_mean  -0.12287   0.04874   0.1382   0.2356   0.40810
## smoothness_mean 133.04711 188.72514 208.3203 240.4691 292.78315
## symmetry_mean  -20.78431  -2.58098   6.6107  15.3513  33.61856
## fractal_dimension_mean -865.96422 -744.19734 -688.0979 -640.5828 -522.87375
## texture_se    -2.98513  -1.91367  -1.4060  -0.9046  -0.02046
## smoothness_se -451.60599 -315.71718 -248.1793 -175.6460 -15.88348
## compactness_se -30.94268   3.70100  23.7214  42.2483  72.24119
## concavity_se   -25.05991 -10.98320  -2.0636   7.6320  23.35489
## concave_points_se 285.73948 381.17421 414.5063 459.1587 565.12707
## symmetry_se    -90.66652 -27.98892   0.3550  29.4079  98.93149
## fractal_dimension_se -735.58558 -423.48834 -261.8735 -89.4792 204.34667
## texture_worst    0.02234   0.18775   0.2620   0.3426   0.52044
## smoothness_worst -55.15123 -18.34344   1.8441  17.3100  53.51026
## symmetry_worst  -2.53864   7.09859  11.6060  17.1774  30.02571
## fractal_dimension_worst 81.74641 131.53188 155.7603 178.0896 233.41108

```

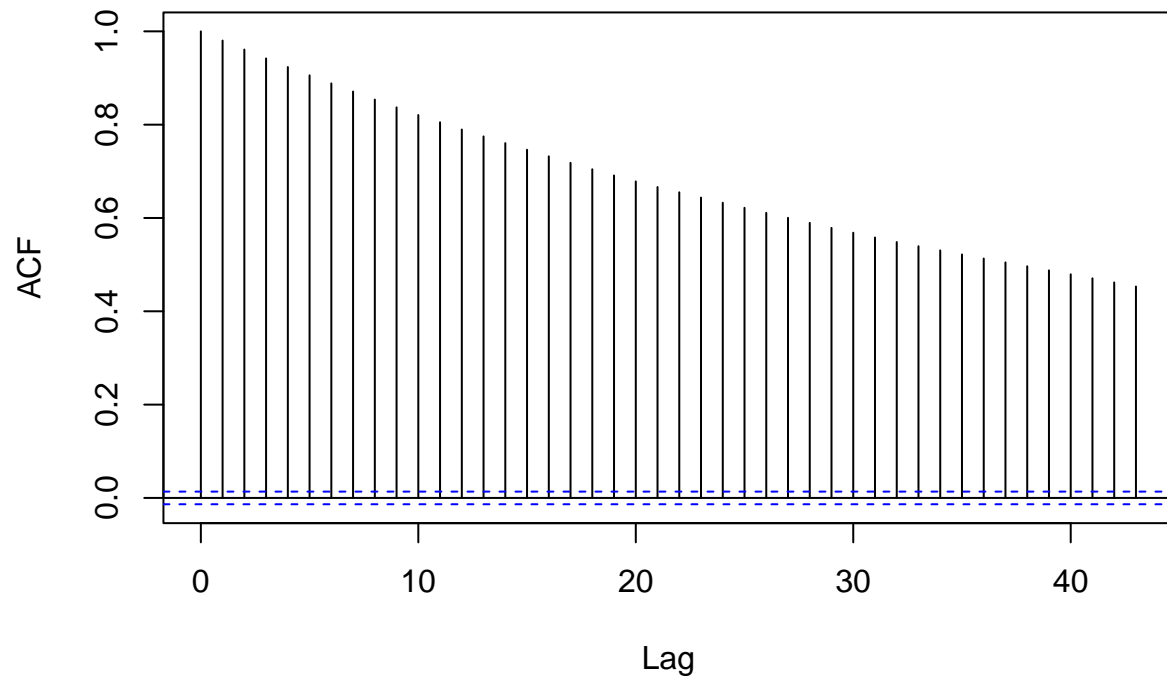
```
acf(out[,1])
```

**Series out[, 1]**



```
acf(out[,2])
```

## Series out[, 2]



*# Correct autocorrelation*

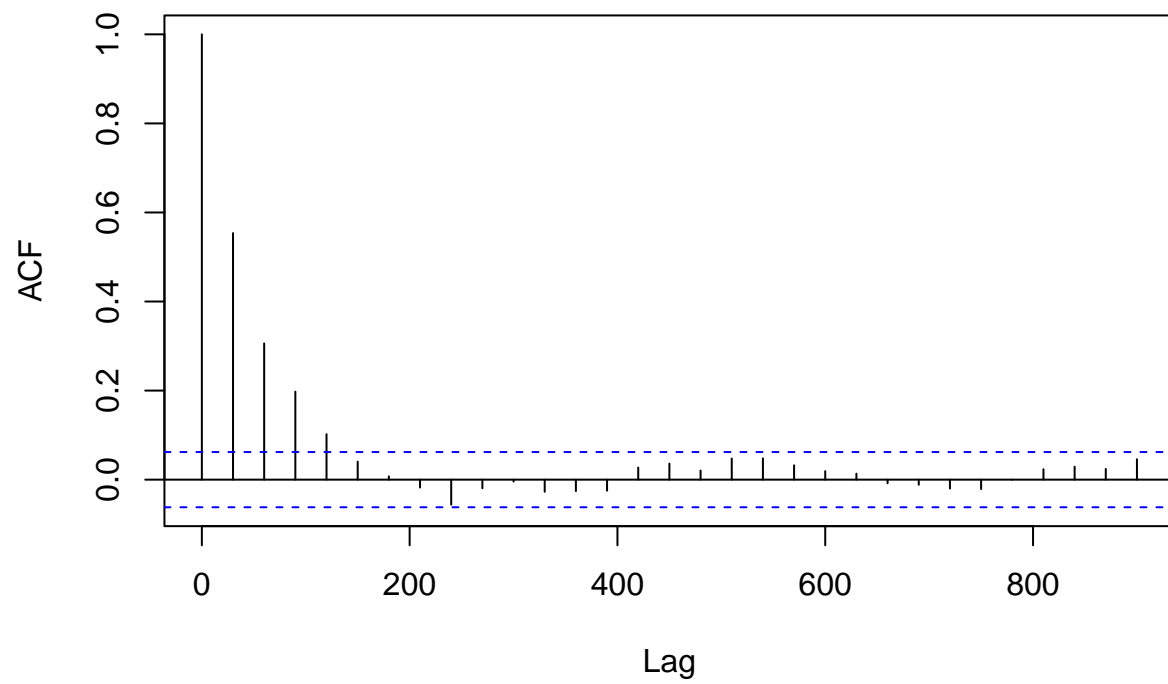
```
out = MCMClogit(formula, data, burnin=5000, mcmc=30000, thin = 30,
                beta.start = beta.start1)
summary(out)
```

```
##
## Iterations = 5001:34971
## Thinning interval = 30
## Number of chains = 1
## Sample size per chain = 1000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##              Mean          SD Naive SE Time-series SE
## (Intercept)    -6.9619    2.9564 0.093488      0.174505
## texture_mean     0.1449     0.1338 0.004230      0.008603
## smoothness_mean 210.3856   38.8092 1.227254      2.364636
## symmetry_mean     6.7972   13.7695 0.435430      0.874829
## fractal_dimension_mean -692.9451 85.9090 2.716683      5.109411
## texture_se       -1.4312     0.7698 0.024345      0.043598
## smoothness_se    -245.3894 120.4758 3.809779      6.431210
## compactness_se    20.8806   28.9179 0.914463      1.881365
## concavity_se      -2.1651   11.8429 0.374504      0.841066
## concave_points_se 419.8216   68.7222 2.173187      4.788362
```

```
## symmetry_se          2.8399  51.6192 1.632341      3.064199
## fractal_dimension_se -235.3213 230.1363 7.277549     12.459032
## texture_worst        0.2585   0.1235 0.003904      0.008201
## smoothness_worst     0.4922  26.7191 0.844933      1.548509
## symmetry_worst       12.2236   8.5569 0.270594      0.595713
## fractal_dimension_worst 153.6350 38.6341 1.221717      2.154905
##
## 2. Quantiles for each variable:
##
##              2.5%      25%      50%      75%      97.5%
## (Intercept)    -1.257e+01 -8.82359 -7.1256 -4.9263 -1.06303
## texture_mean   -1.264e-01  0.05609  0.1493  0.2347  0.42288
## smoothness_mean 1.346e+02 185.56985 207.4429 237.4691 285.25906
## symmetry_mean  -2.133e+01 -2.07081  7.2447 16.8415 33.46638
## fractal_dimension_mean -8.685e+02 -748.06764 -687.5111 -638.0705 -535.08845
## texture_se     -3.011e+00 -1.90562 -1.3896 -0.9116 -0.01158
## smoothness_se  -4.695e+02 -329.76049 -253.4453 -166.7919  2.89369
## compactness_se -3.356e+01  1.54953  21.3018 40.3877 76.28100
## concavity_se   -2.669e+01 -10.76965 -2.0245  6.0901 20.70983
## concave_points_se 2.861e+02 374.51992 416.7823 463.6432 566.28489
## symmetry_se    -1.014e+02 -29.30766  2.0131 37.1225 99.63610
## fractal_dimension_se -7.050e+02 -385.97951 -235.5333 -66.6496 194.69516
## texture_worst   2.987e-03  0.17517  0.2570  0.3330  0.52035
## smoothness_worst -5.199e+01 -17.05354 -0.2086 17.8640 54.16787
## symmetry_worst  -2.732e+00  6.80055 11.5672 17.4516 30.08029
## fractal_dimension_worst 7.957e+01 126.52236 155.9399 178.0914 231.92666
```

```
acf(out[,1])
```

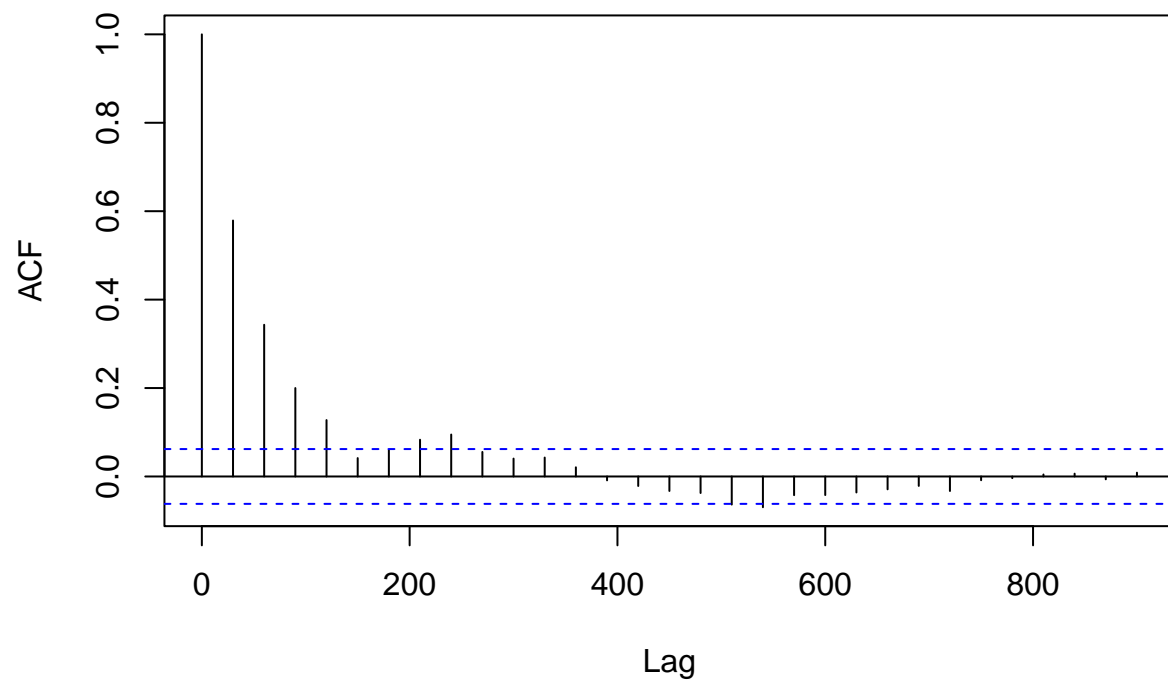
**Series out[, 1]**



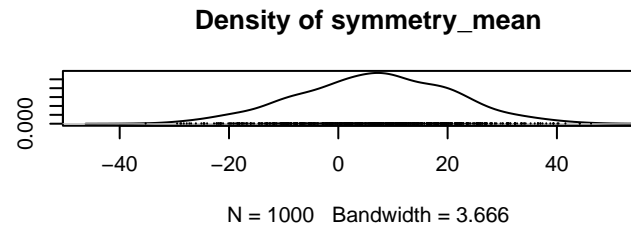
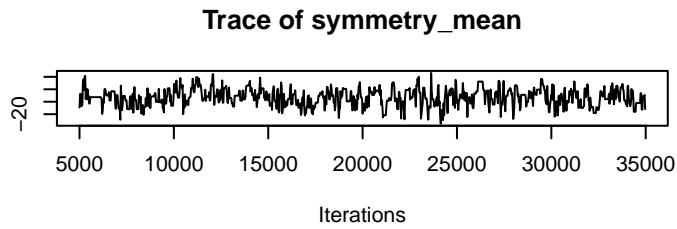
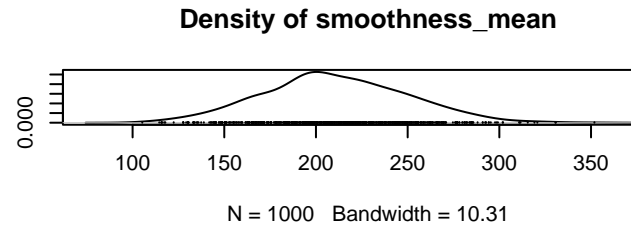
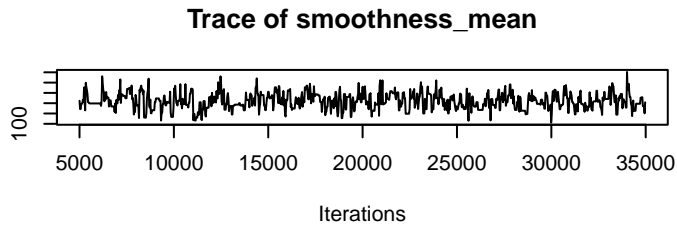
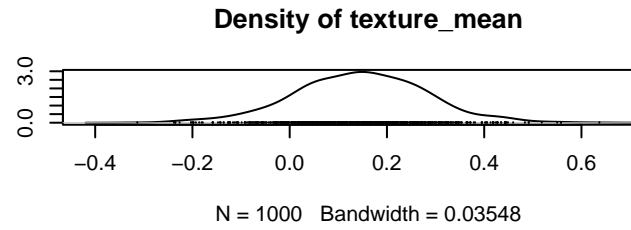
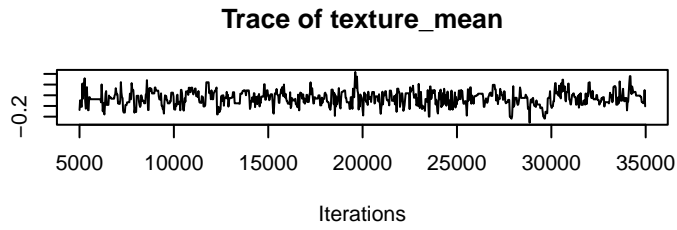
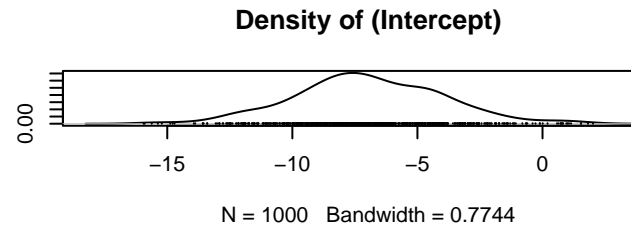
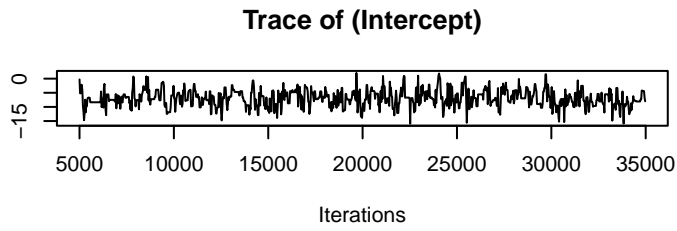
```
acf(out[,2])
```



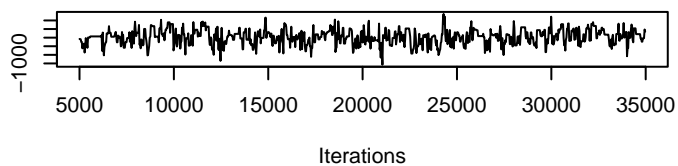
**Series out[, 2]**



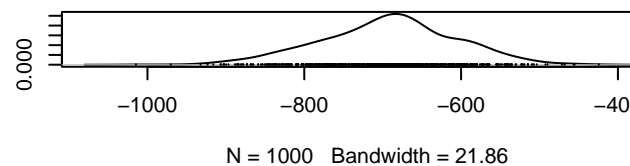
```
plot(out)
```



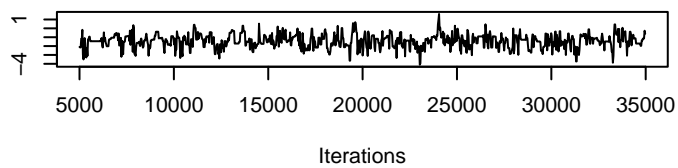
**Trace of fractal\_dimension\_mean**



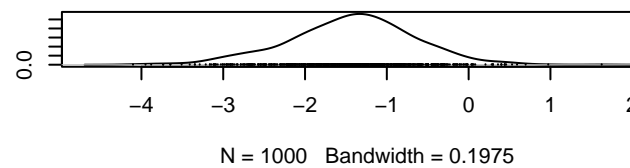
**Density of fractal\_dimension\_mean**



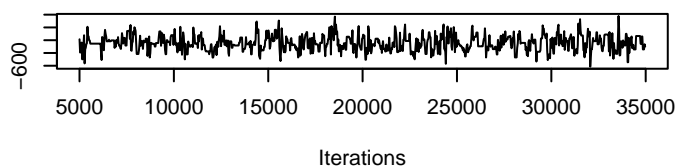
**Trace of texture\_se**



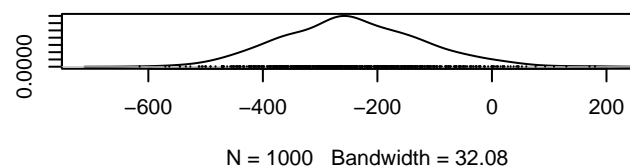
**Density of texture\_se**



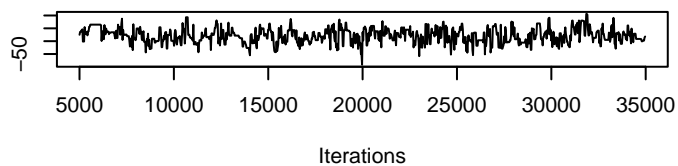
**Trace of smoothness\_se**



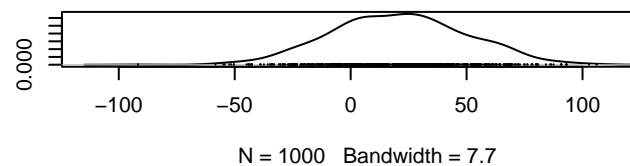
**Density of smoothness\_se**



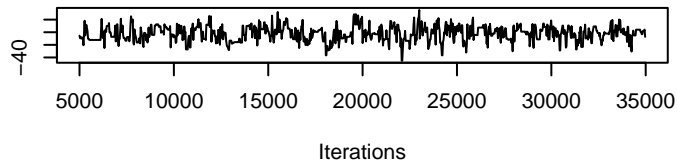
**Trace of compactness\_se**



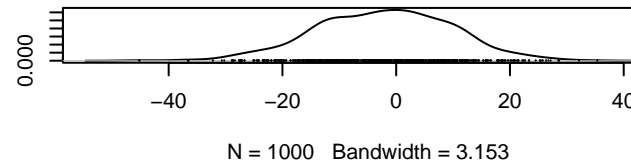
**Density of compactness\_se**



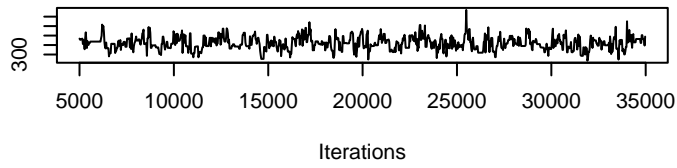
**Trace of concavity\_se**



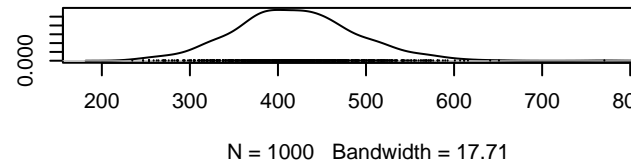
**Density of concavity\_se**



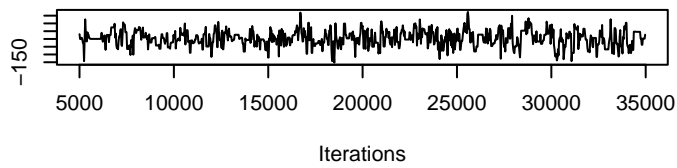
**Trace of concave\_points\_se**



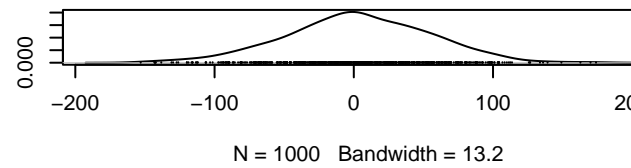
**Density of concave\_points\_se**



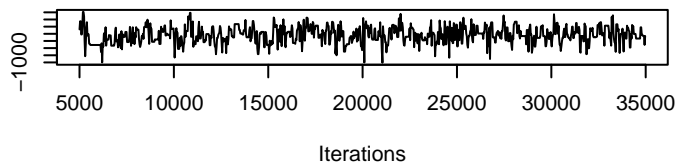
**Trace of symmetry\_se**



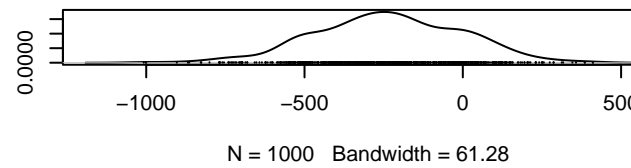
**Density of symmetry\_se**

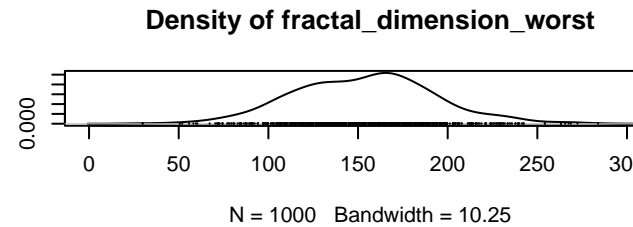
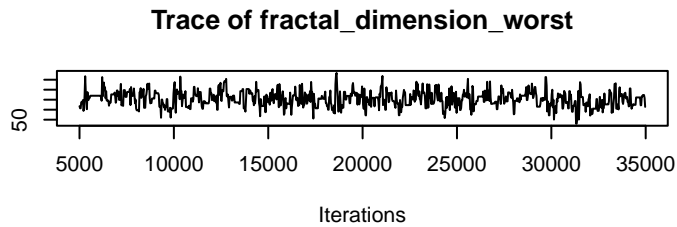
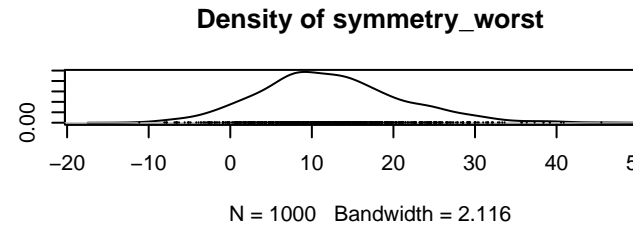
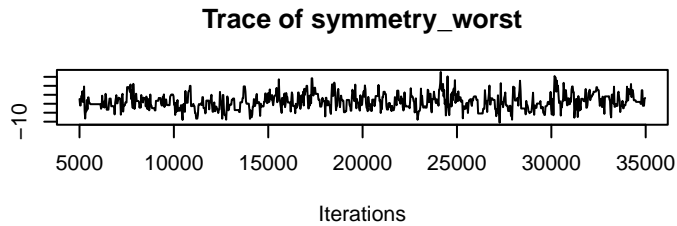
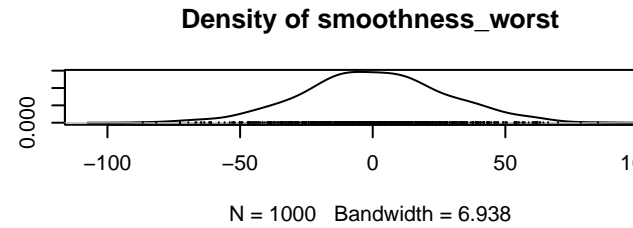
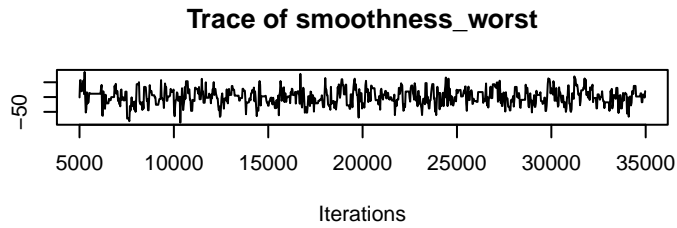
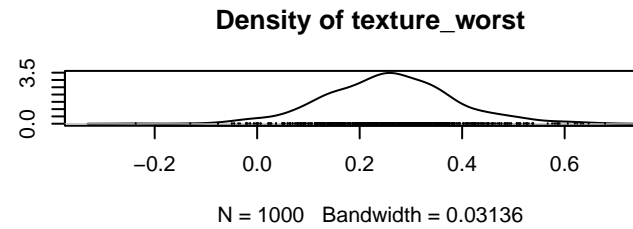
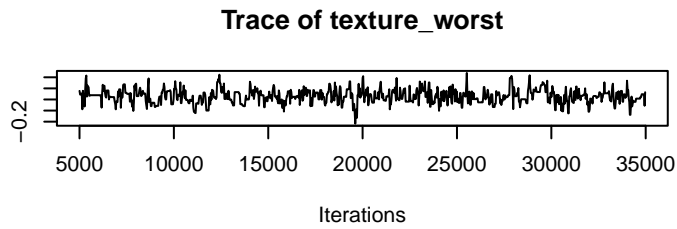


**Trace of fractal\_dimension\_se**



**Density of fractal\_dimension\_se**





## Bayes var selection

```
formula_str_b <- paste("diagnosis ~", paste(selected_bayes, collapse = " + "))
formula_b <- as.formula(formula_str_b)
```

```
# starting point
freq_model <- lm(formula_b, data = data)
beta.start <- coef(freq_model)
```

```
out_b = MCMClogit(formula_b, data, burnin=1000, mcmc=21000)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(out_b)
```

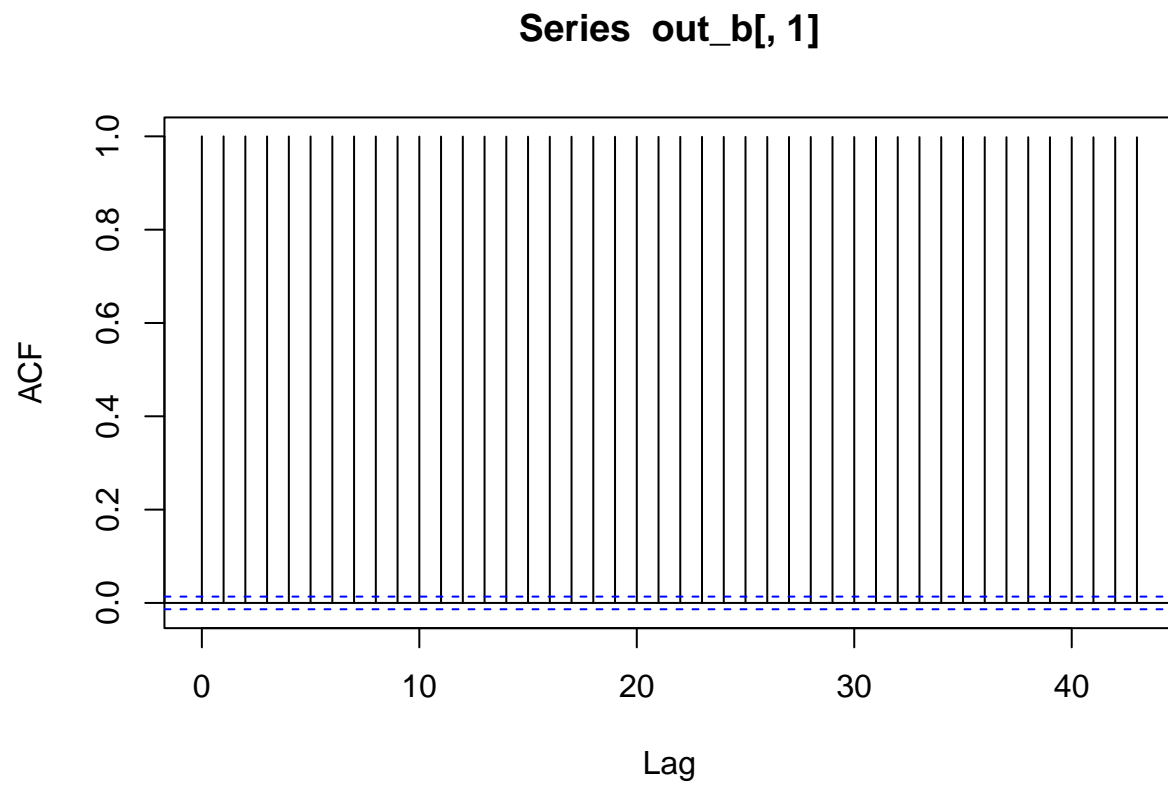
```
##
```

```

## Iterations = 1001:22000
## Thinning interval = 1
## Number of chains = 1
## Sample size per chain = 21000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##
##              Mean SD Naive SE Time-series SE
## (Intercept)      -66.8917  0      0      0
## perimeter_mean    -0.3481  0      0      0
## concave_points_mean 113.2543  0      0      0
## compactness_mean   -67.0638  0      0      0
## concavity_mean     36.5866  0      0      0
## area_se           0.1428  0      0      0
## smoothness_se     433.2403  0      0      0
## concave_points_se  428.0355  0      0      0
## fractal_dimension_se -2564.0379  0      0      0
## radius_worst       3.4530  0      0      0
## texture_worst      0.4372  0      0      0
## fractal_dimension_worst 318.5212  0      0      0
##
## 2. Quantiles for each variable:
##
##              2.5%      25%      50%      75%      97.5%
## (Intercept)      -66.8917 -66.8917 -66.8917 -66.8917 -66.8917
## perimeter_mean    -0.3481 -0.3481 -0.3481 -0.3481 -0.3481
## concave_points_mean 113.2543 113.2543 113.2543 113.2543 113.2543
## compactness_mean   -67.0638 -67.0638 -67.0638 -67.0638 -67.0638
## concavity_mean     36.5866  36.5866  36.5866  36.5866  36.5866
## area_se           0.1428  0.1428  0.1428  0.1428  0.1428
## smoothness_se     433.2403 433.2403 433.2403 433.2403 433.2403
## concave_points_se  428.0355 428.0355 428.0355 428.0355 428.0355
## fractal_dimension_se -2564.0379 -2564.0379 -2564.0379 -2564.0379 -2564.0379
## radius_worst       3.4530  3.4530  3.4530  3.4530  3.4530
## texture_worst      0.4372  0.4372  0.4372  0.4372  0.4372
## fractal_dimension_worst 318.5212 318.5212 318.5212 318.5212 318.5212

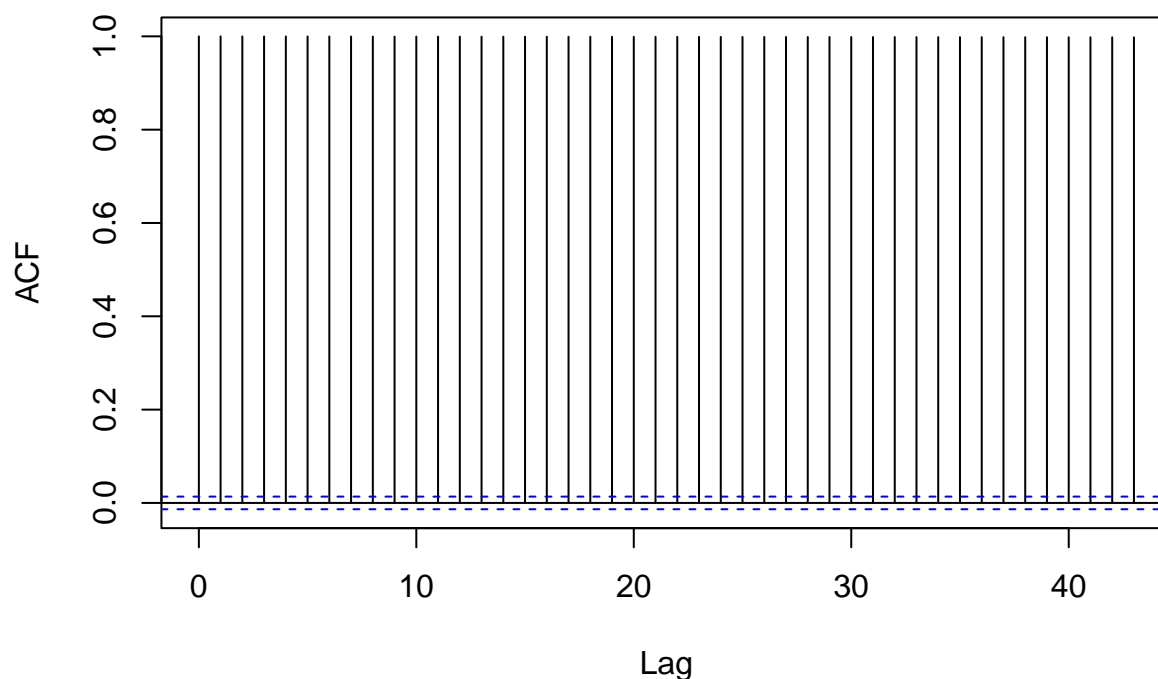
```

`acf(out_b[,1])`



```
acf(out_b[,2])
```

## Series out\_b[, 2]



```
out_b = MCMClogit(formula_b, data, burnin=5000, mcmc=50000,
                  beta.start = beta.start, thin = 50, tune=0.5)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
summary(out_b)
```

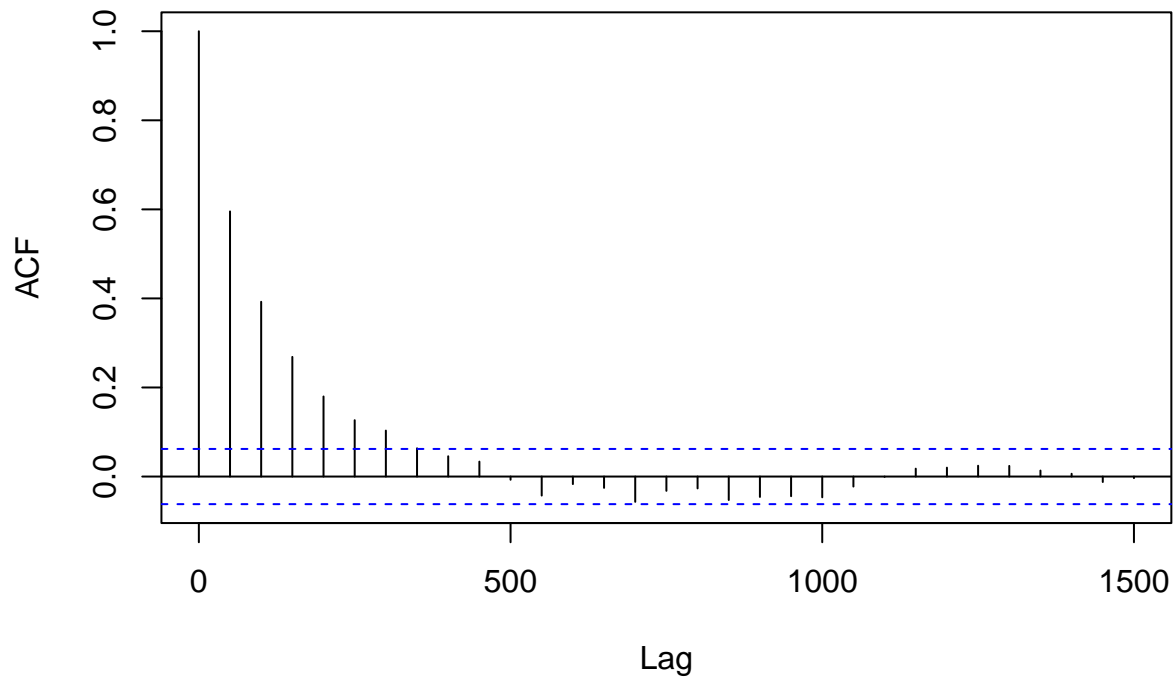
```
##
## Iterations = 5001:54951
## Thinning interval = 50
## Number of chains = 1
## Sample size per chain = 1000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##              Mean          SD Naive SE Time-series SE
## (Intercept)  -4.089e+01    6.10325 1.930e-01    4.069e-01
## perimeter_mean -8.767e-02    0.08165 2.582e-03    5.462e-03
## concave_points_mean  7.789e+01 33.64842 1.064e+00    2.578e+00
## compactness_mean  -4.411e+01 20.79019 6.574e-01    1.713e+00
## concavity_mean   1.849e+01 10.76886 3.405e-01    8.339e-01
## area_se        -7.147e-03    0.01074 3.396e-04    5.185e-04
## smoothness_se    2.135e+02 150.05698 4.745e+00    1.353e+01
## concave_points_se  2.921e+02 101.48434 3.209e+00    5.821e+00
## fractal_dimension_se -1.196e+03 340.24976 1.076e+01    1.952e+01
## radius_worst     1.662e+00    0.44882 1.419e-02    2.645e-02
```



```
## texture_worst          2.749e-01    0.05205 1.646e-03      3.775e-03
## fractal_dimension_worst 1.559e+02  46.60596 1.474e+00      3.093e+00
##
## 2. Quantiles for each variable:
##
##              2.5%      25%      50%      75%      97.5%
## (Intercept)    -5.348e+01 -4.504e+01 -4.084e+01 -3.637e+01 -29.41492
## perimeter_mean -2.363e-01 -1.449e-01 -9.337e-02 -3.226e-02  0.08183
## concave_points_mean 1.352e+01 5.524e+01 7.861e+01 1.017e+02 141.59098
## compactness_mean -8.455e+01 -5.933e+01 -4.349e+01 -2.957e+01 -4.49315
## concavity_mean -2.063e+00 1.172e+01 1.872e+01 2.589e+01 39.71891
## area_se        -2.774e-02 -1.452e-02 -6.903e-03 -1.231e-04  0.01324
## smoothness_se  -8.381e+01 1.137e+02 2.063e+02 3.125e+02 512.29411
## concave_points_se 8.751e+01 2.252e+02 2.941e+02 3.573e+02 499.56987
## fractal_dimension_se -1.905e+03 -1.411e+03 -1.177e+03 -9.488e+02 -589.19519
## radius_worst    7.393e-01 1.375e+00 1.711e+00 1.994e+00 2.41334
## texture_worst    1.752e-01 2.396e-01 2.744e-01 3.086e-01 0.37594
## fractal_dimension_worst 7.238e+01 1.237e+02 1.556e+02 1.866e+02 255.52220
```

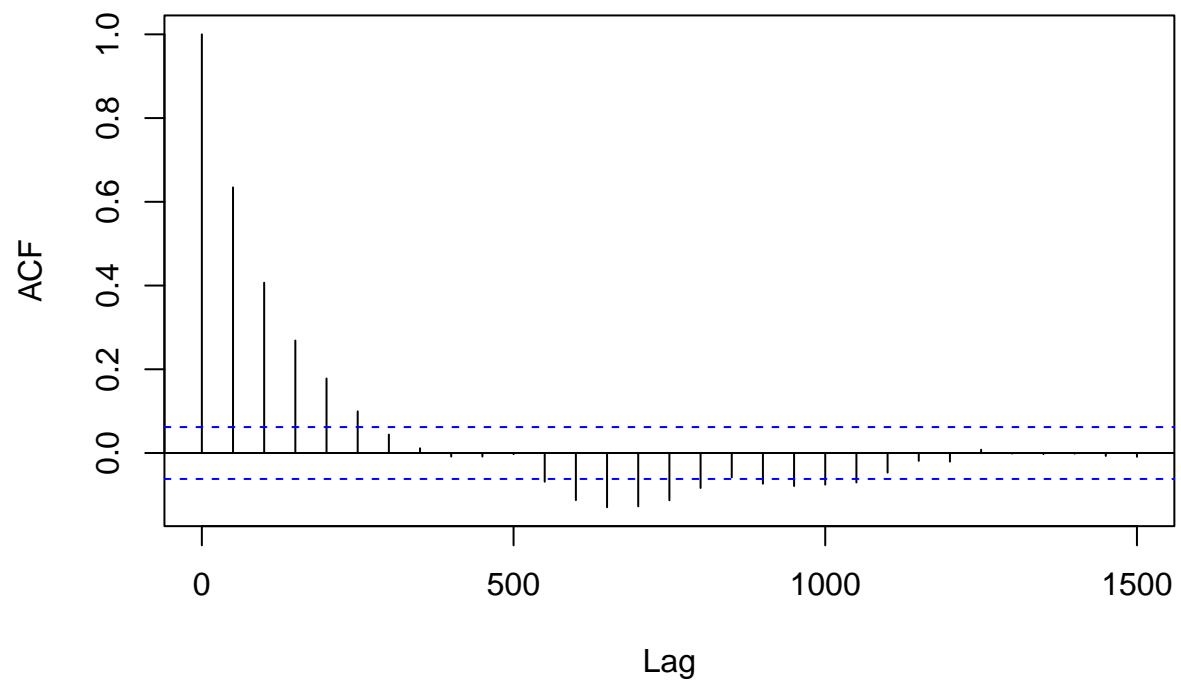
```
acf(out_b[,1])
```

### Series out\_b[, 1]



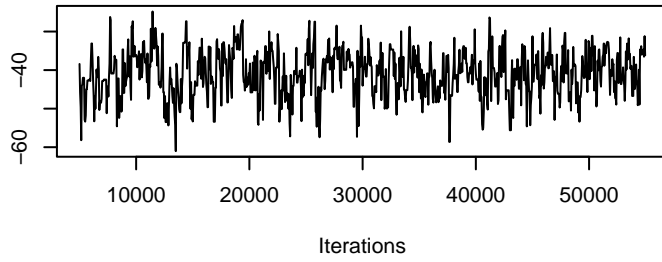
```
acf(out_b[,2])
```

**Series out\_b[, 2]**

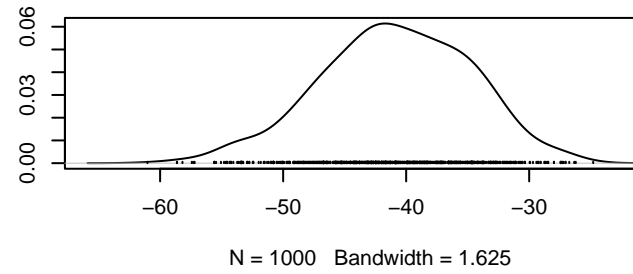


```
plot(out_b)
```

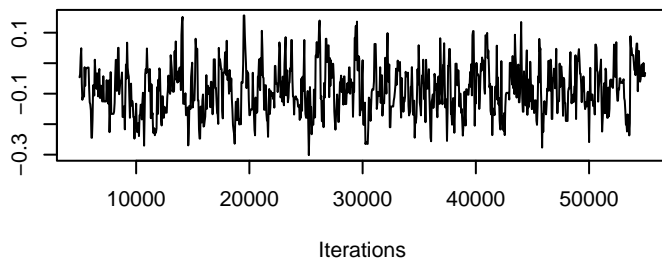
**Trace of (Intercept)**



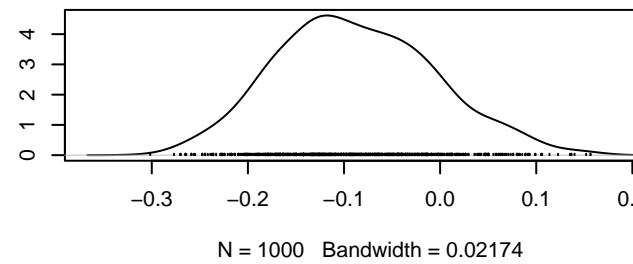
**Density of (Intercept)**



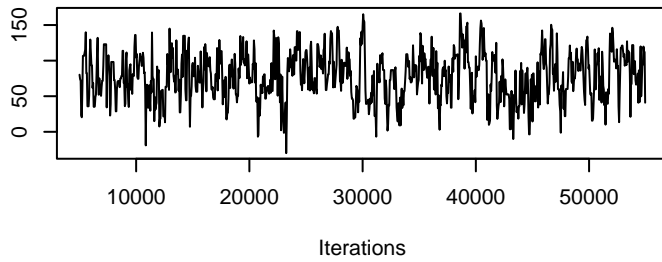
**Trace of perimeter\_mean**



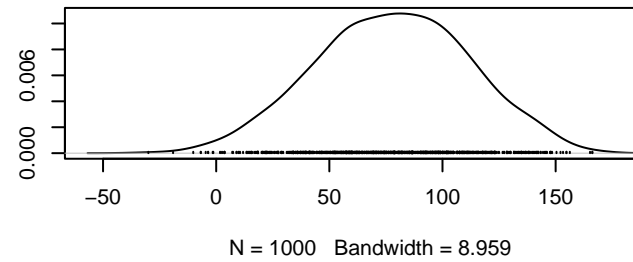
**Density of perimeter\_mean**



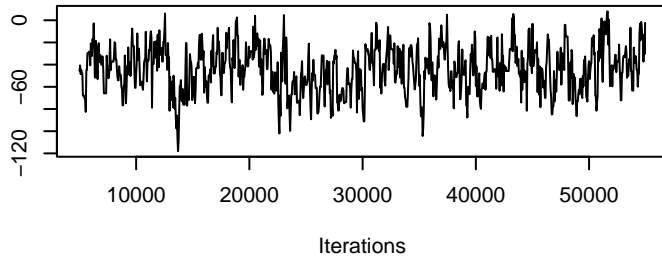
**Trace of concave\_points\_mean**



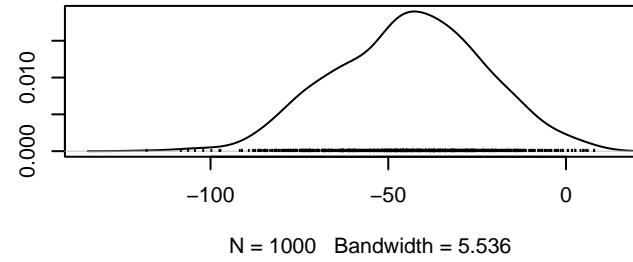
**Density of concave\_points\_mean**



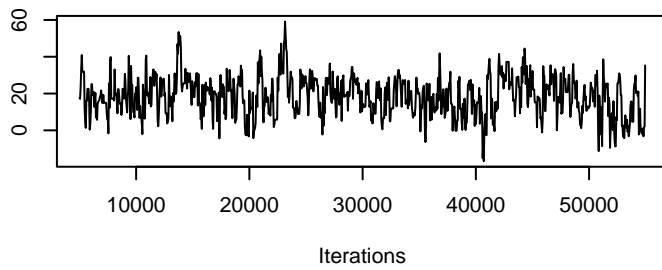
**Trace of compactness\_mean**



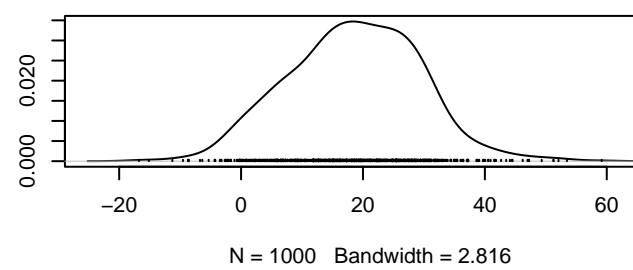
**Density of compactness\_mean**



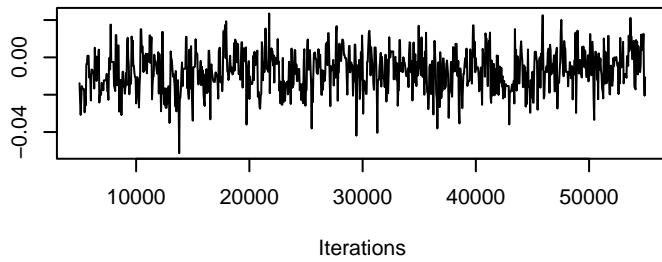
**Trace of concavity\_mean**



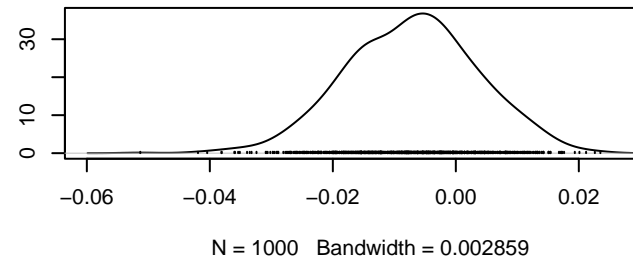
**Density of concavity\_mean**

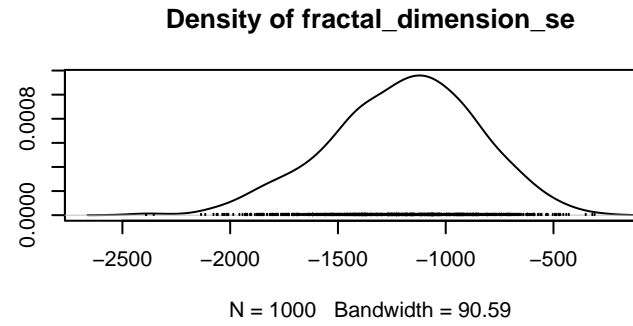
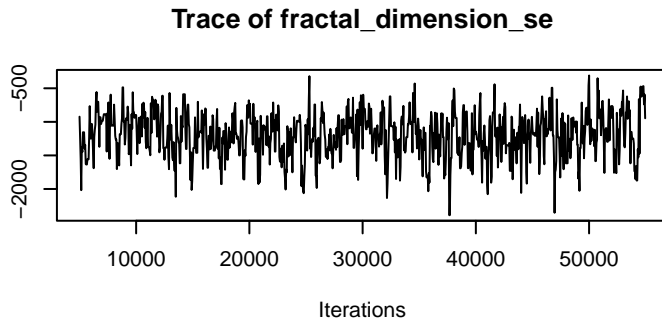
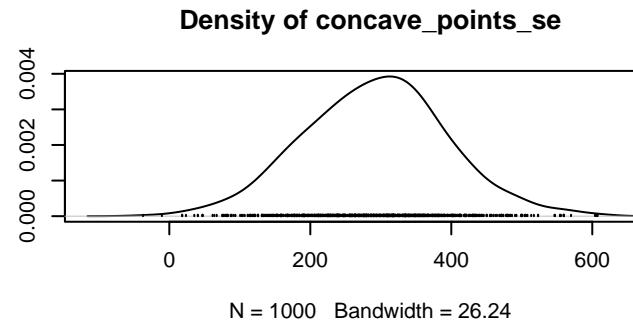
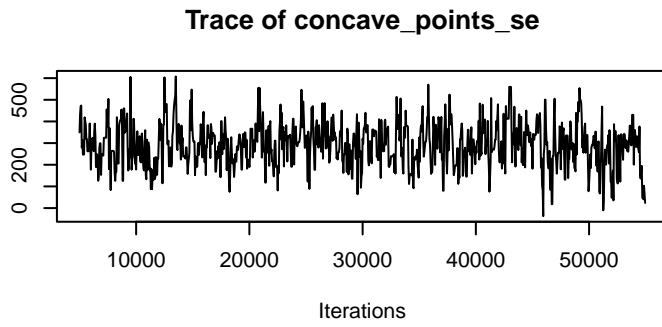
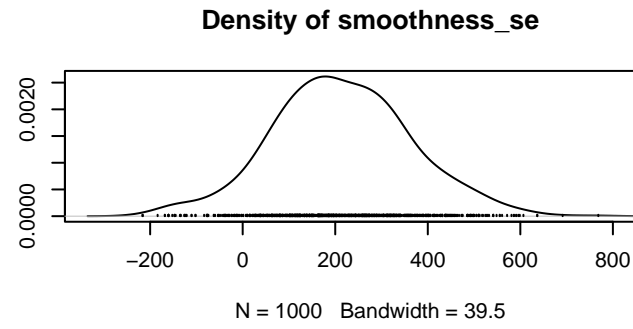
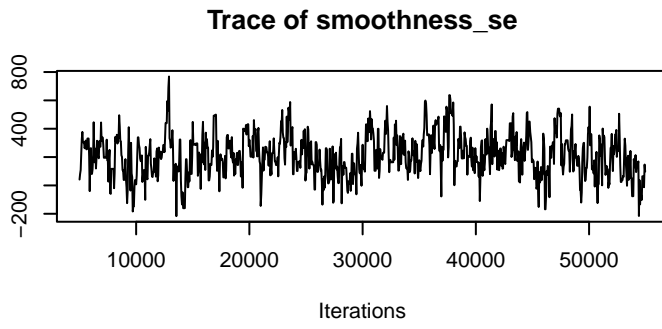


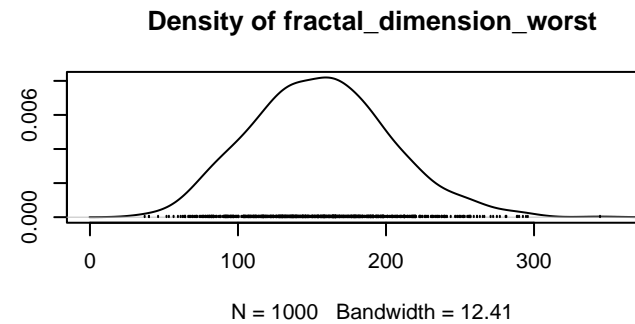
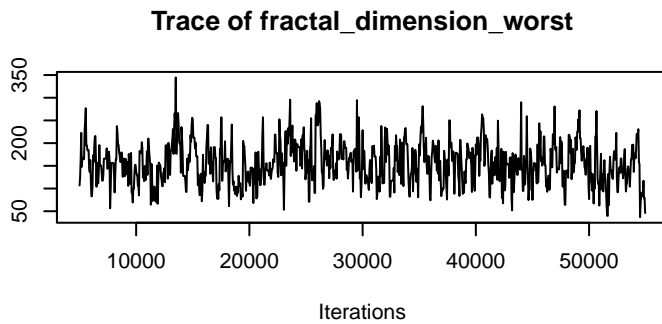
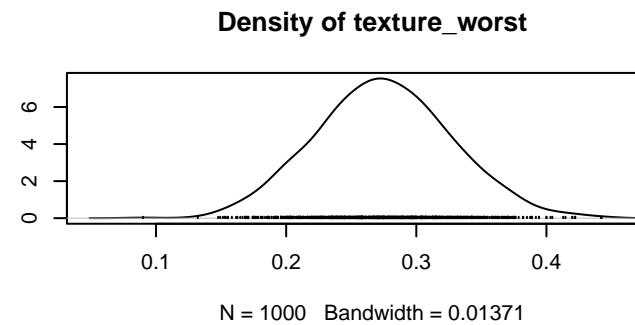
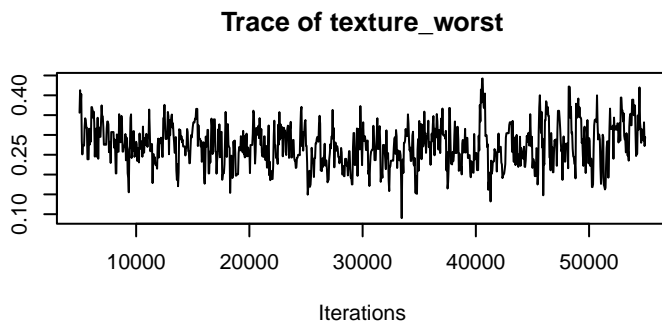
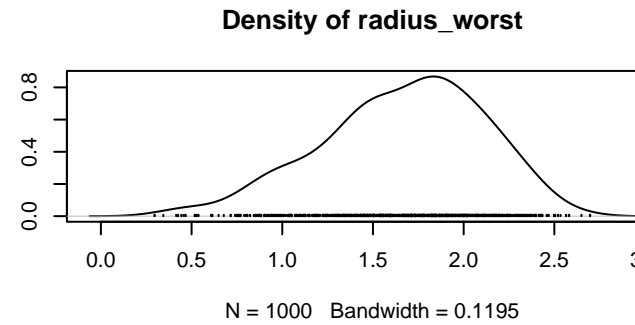
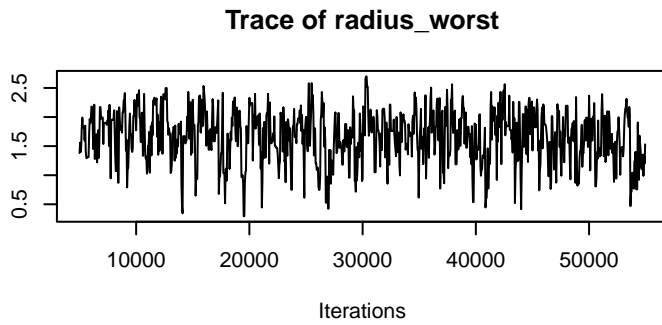
**Trace of area\_se**



**Density of area\_se**







```
# Run Bayesian logistic regression with Spike-and-Slab priors (to variable selection)
set.seed(123)
model <- MCMClogit(formula, data = data, burnin = 5000, mcmc = 20000
, marginal.likelihood = "Laplace") # , thin = 10, b0 = 0, B0 = 0.1
```

```
## Warning in MCMClogit(formula, data = data, burnin = 5000, mcmc = 20000, : Cannot calculate marginal likelihood
summary(model)
```

```
##
## Iterations = 5001:25000
## Thinning interval = 1
## Number of chains = 1
## Sample size per chain = 20000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
```

```
##              Mean      SD Naive SE Time-series SE
## (Intercept)    -6.8170   2.9513 0.0208690      0.219795
## texture_mean     0.1294   0.1348 0.0009531      0.009661
## smoothness_mean  207.3107 41.2005 0.2913315      3.233150
## symmetry_mean     6.9549 13.6898 0.0968017      0.997838
## fractal_dimension_mean -684.5427 91.4640 0.6467481      7.597434
## texture_se       -1.4813   0.7652 0.0054108      0.054139
## smoothness_se    -234.4591 126.4486 0.8941266      9.664144
## compactness_se    21.2337  26.1173 0.1846775      1.799246
## concavity_se      -1.4548 11.8178 0.0835642      0.846533
## concave_points_se 411.5040 64.7565 0.4578975      4.564841
## symmetry_se       -1.2348 53.9843 0.3817264      3.947217
## fractal_dimension_se -193.6926 218.7169 1.5465618     12.937977
## texture_worst      0.2698   0.1255 0.0008877      0.009101
## smoothness_worst   0.9483  28.4640 0.2012712      2.053050
## symmetry_worst     12.2825   8.7408 0.0618070      0.658451
## fractal_dimension_worst 147.8957 37.6905 0.2665122      2.554404
##
## 2. Quantiles for each variable:
##
##              2.5%      25%      50%      75%      97.5%
## (Intercept)   -12.51578  -8.78930  -6.8489  -4.7347  -0.98681
## texture_mean   -0.14451   0.03281   0.1240   0.2273   0.38854
## smoothness_mean 126.36954 181.04922 207.4663 233.5748 289.55512
## symmetry_mean  -20.77997  -1.01465   7.3981  15.9900  33.16468
## fractal_dimension_mean -876.85892 -744.76209 -676.1047 -624.8130 -519.74380
## texture_se     -2.98677  -1.92703  -1.5112  -0.9390  -0.04385
## smoothness_se  -492.69637 -319.84949 -241.2256 -153.4581  38.46158
## compactness_se  -30.59198   5.34563  19.9356  38.9766  73.12934
## concavity_se    -25.61537  -8.41350  -1.4880   6.0837  21.32527
## concave_points_se 296.11128 368.03285 407.6404 450.6254 546.85401
## symmetry_se    -105.48846  -39.57962   1.0433  37.1174 108.12212
## fractal_dimension_se -656.33024 -324.37467 -186.6719 -45.9750 200.42144
## texture_worst    0.04377   0.17585   0.2714   0.3511   0.53608
## smoothness_worst -54.51081 -18.36763   0.8286  21.8650  56.01089
## symmetry_worst   -4.58397   6.13773  12.0131  18.5500  29.77524
## fractal_dimension_worst 82.79661 124.27778 144.7862 172.0666 233.27842
##
## apply(model, 2, function(x) mean(x != 0)) # Approximate inclusion probability
##
##              (Intercept)          texture_mean          smoothness_mean
##              1              1              1
##              symmetry_mean fractal_dimension_mean          texture_se
##              1              1              1
##              smoothness_se          compactness_se          concavity_se
##              1              1              1
##              concave_points_se          symmetry_se fractal_dimension_se
##              1              1              1
##              texture_worst          smoothness_worst          symmetry_worst
##              1              1              1
## fractal_dimension_worst
##              1
```

## Evaluate Models with Deviance Information Criterion (DIC)

In the following code, we will calculate the Deviance Information Criterion (DIC) for both the Frequentist and Bayesian models. The DIC is a measure of model fit that penalizes the complexity of the model. Lower values of DIC indicate better model fit. The DIC is calculated as follows:

$$DIC = \bar{D} + p_D$$

where:

-  $\bar{D}$  is the posterior mean deviance:

$$\bar{D} = \mathbb{E}[D(\theta) \mid \mathcal{D}]$$

with  $D(\theta) = -2\log p(\mathcal{D} \mid \theta)$ , the deviance evaluated at parameter  $\theta$ . -  $p_D$  a penalization term (effective number of parameters to penalize model complexity):

$$p_D = \bar{D} - D(\hat{\theta})$$

where  $\hat{\theta}$  is the posterior mean of  $\theta$ .

The R implementation of the DIC function is as follows and was developed with help of Prof Michael Wiper:

```
# DIC Code
DIC = function(model, X, data, target) {
  dev = 0
  # Calculate Average Deviance of MCMC
  for (i in 1:nrow(model)) {
    params <- model[i,]
    p = inv.logit(X %*% params)
    p[data[target] == 0] = 1-p[data[target] == 0]
    dev = dev - 2 * sum(log(p)) # Negative log-likelihood
  }
  D_bar = dev / nrow(model)

  # D_theta: Deviance at the posterior mean (using the average parameter values)
  posterior_means <- colMeans(model)
  linear_predictor <- X %*% posterior_means
  p_post <- inv.logit(linear_predictor)
  p_post[data[target] == 0] = 1-p_post[data[target] == 0]

  D_theta = -2 * sum(log(p_post)) # Deviance at the posterior mean

  # p_D: Posterior deviance penalty
  p_D = D_bar - D_theta

  # DIC
  DIC = D_bar + p_D

  return(list(DIC=DIC, D_bar=D_bar, p_D=p_D))
}
```

We now continue with applying the DIC Score to the model derived from frequentist variable selection and the model derived from Bayesian variable selection. The straight forward conclusion is that the DIC is significantly better (lower) for the model that was set up with the Bayesian Variable Selection approach. Based on this result, we conclude this to be the best model and will use it for further analysis.



```

# Frequentist
model = out
X <- model.matrix(~ texture_mean + smoothness_mean + symmetry_mean +
  fractal_dimension_mean + texture_se + smoothness_se + compactness_se +
  concavity_se + concave_points_se + symmetry_se + fractal_dimension_se +
  texture_worst + smoothness_worst + symmetry_worst + fractal_dimension_worst, data = data) # model matrix
target = "diagnosis"

print("Frequentist Variable Selection DIC Score")

## [1] "Frequentist Variable Selection DIC Score"
DIC(model, X, data, target)

## $DIC
## [1] 245.5246
##
## $D_bar
## [1] 229.9737
##
## $p_D
## [1] 15.5509

# Bayesian
model = out_b
X <- model.matrix(~ perimeter_mean + concave_points_mean + compactness_mean +
  concavity_mean + area_se + smoothness_se + concave_points_se +
  fractal_dimension_se + radius_worst + texture_worst + fractal_dimension_worst, data = data) # model matrix
target = "diagnosis"

print("Bayesian Variable Selection DIC Score")

## [1] "Bayesian Variable Selection DIC Score"
DIC(model, X, data, target)

## $DIC
## [1] 89.86894
##
## $D_bar
## [1] 80.33019
##
## $p_D
## [1] 9.538756

```

## Prediction

```
# TODO: Make predictions
```

## Conclusion

TODO: Write Conclusion